

ICS 143 - Principles of Operating Systems

Lectures 3 and 4 - Processes and Threads

Prof. Nalini Venkatasubramanian

nalini@ics.uci.edu

Outline

- Process Concept
 - Process Scheduling
 - Operations on Processes
 - Cooperating Processes
 - Threads
 - Interprocess Communication
-

Process Concept

- An operating system executes a variety of programs
 - batch systems - jobs
 - time-shared systems - user programs or tasks
 - job and program used interchangeably
 - Process - a program in execution
 - process execution proceeds in a sequential fashion
 - A process contains
 - program counter, stack and data section
-

Process =? Program

```
main ()
{
    ...;
}


A () {
    ...
}
```

Program

```
main ()
{
    ...;
}

A () {
    ...
}
```

Process

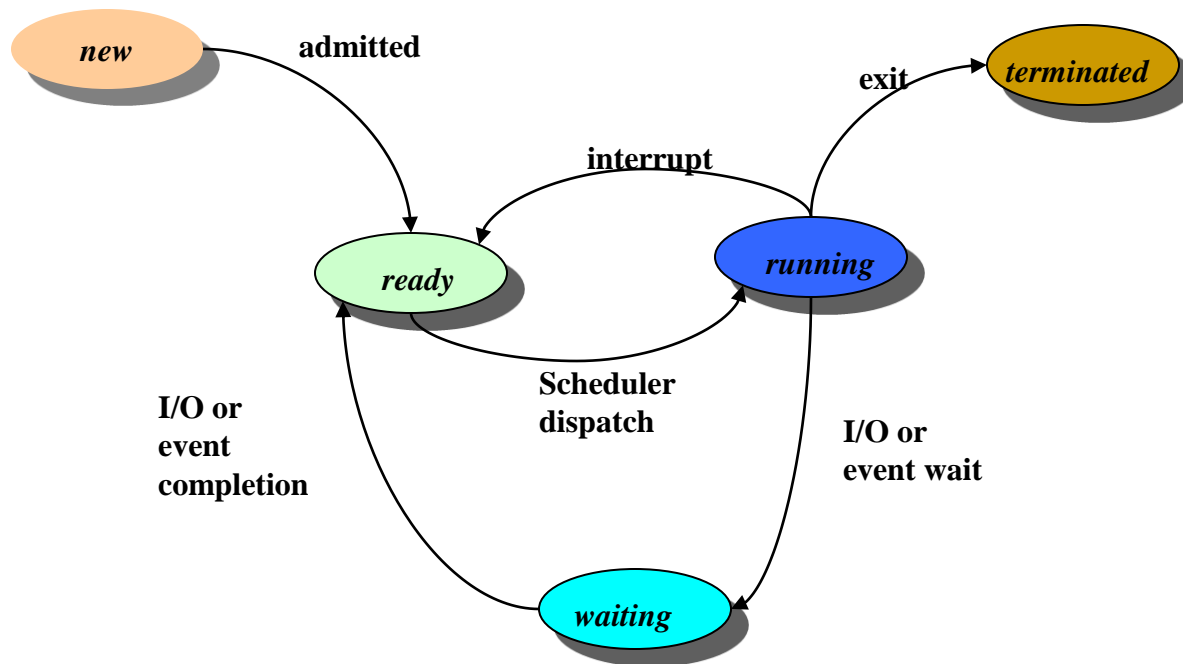


The diagram shows the memory layout of a process. It consists of two stacked rectangular boxes. The top box is labeled 'Heap'. The bottom box is labeled 'Stack' and contains the text 'A' and 'main' stacked vertically, representing the call stack.

- More to a process than just a program:
 - Program is just part of the process state
 - I run emacs on lectures.txt, you run it on homework.java – Same program, different processes
- Less to a process than a program:
 - A program can invoke more than one process
 - cc starts up cpp, cc1, cc2, as, and ld

Process State

- A process changes state as it executes.

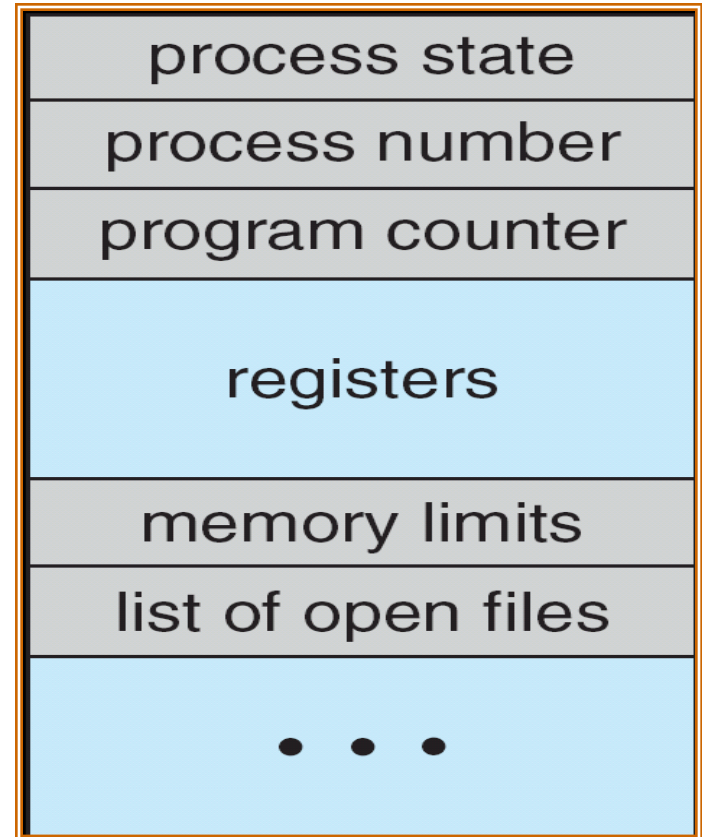


Process States

- New - The process is being created.
 - Running - Instructions are being executed.
 - Waiting - Waiting for some event to occur.
 - Ready - Waiting to be assigned to a processor.
 - Terminated - Process has finished execution.
-

Process Control Block

- Contains information associated with each process
 - Process State - e.g. new, ready, running etc.
 - Process Number – Process ID
 - Program Counter - address of next instruction to be executed
 - CPU registers - general purpose registers, stack pointer etc.
 - CPU scheduling information - process priority, pointer
 - Memory Management information - base/limit information
 - Accounting information - time limits, process number
 - I/O Status information - list of I/O devices allocated

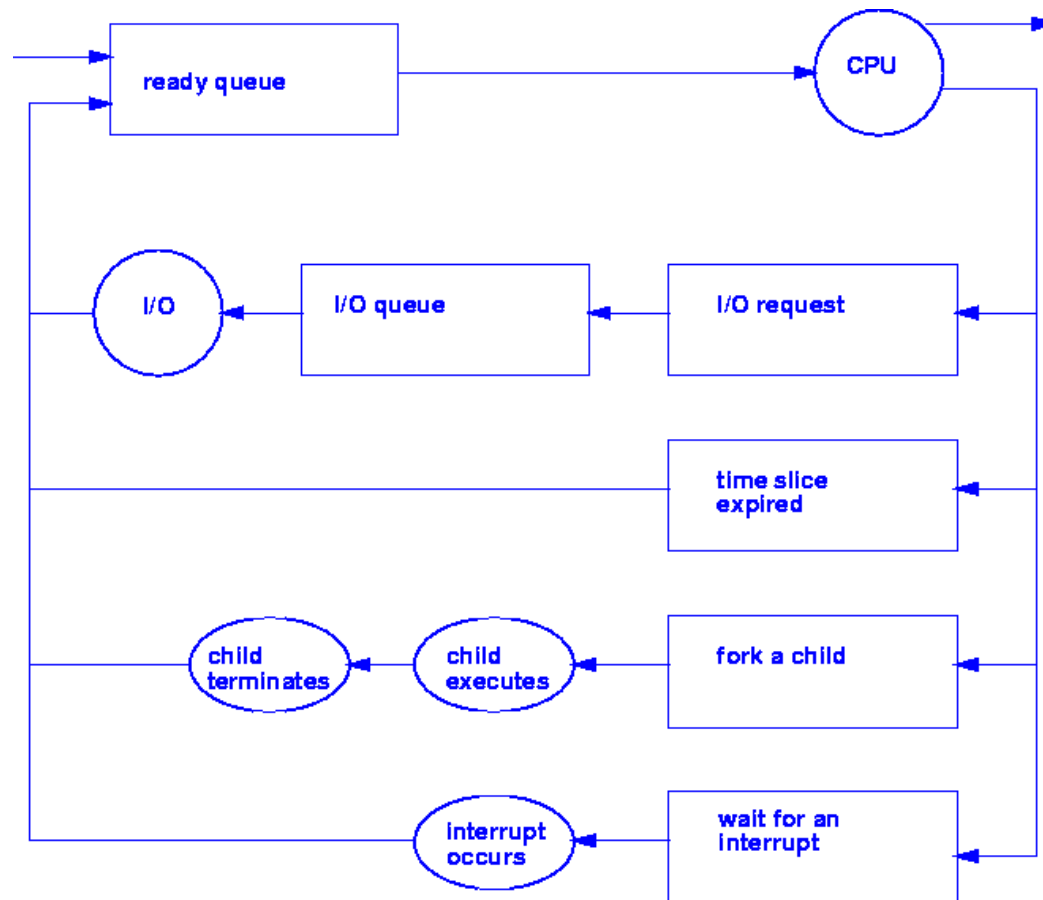


Process
Control
Block

Process Scheduling

Process (PCB) moves from queue to queue

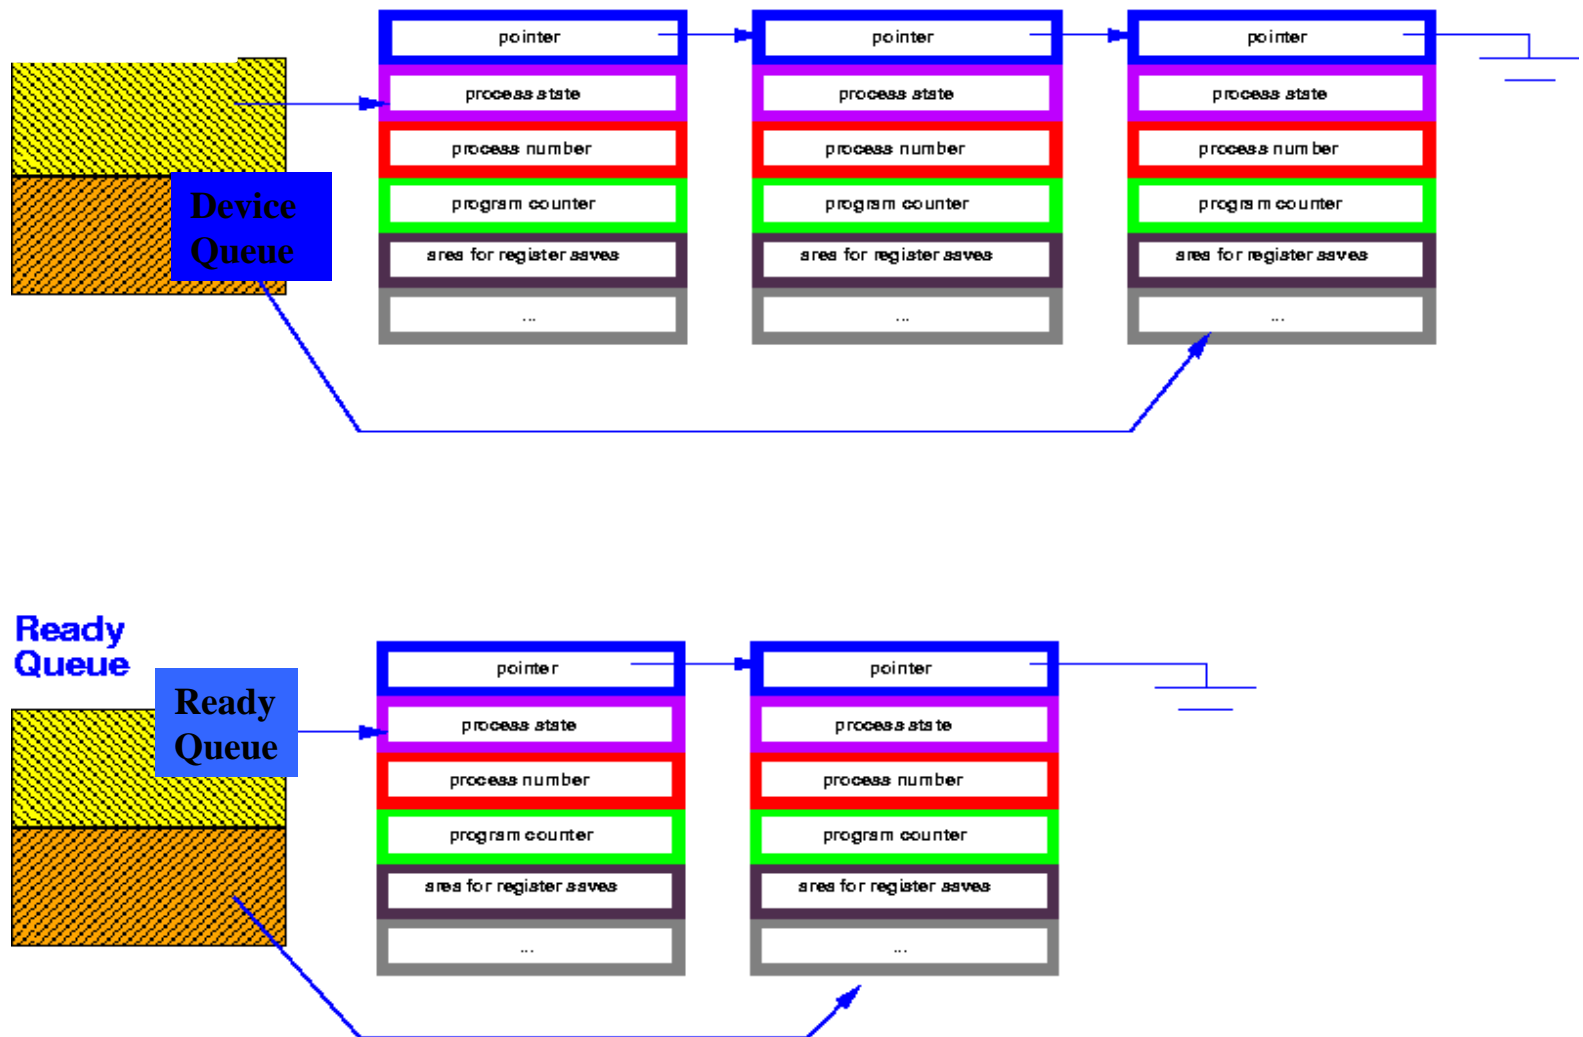
When does it move? Where? A scheduling decision



Process Scheduling Queues

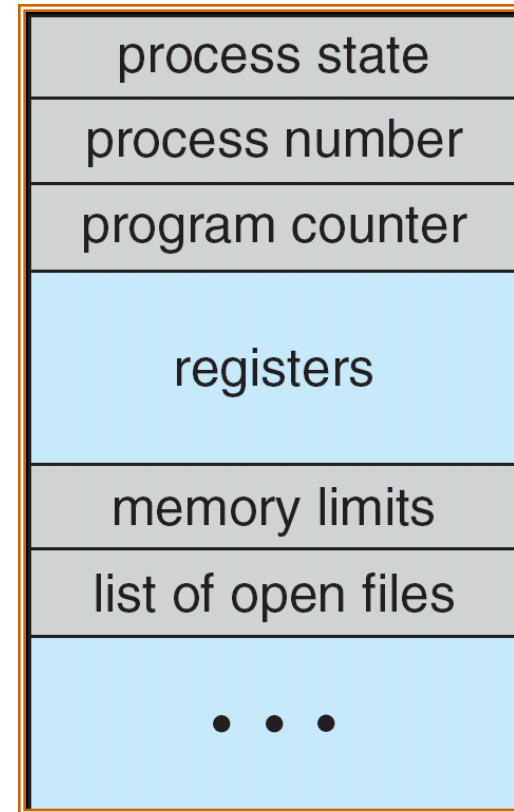
- Job Queue - set of all processes in the system
 - Ready Queue - set of all processes residing in main memory, ready and waiting to execute.
 - Device Queues - set of processes waiting for an I/O device.
 - Process migration between the various queues.
 - Queue Structures - typically linked list, circular list etc.
-

Process Queues



Enabling Concurrency and Protection: Multiplex processes

- Only one process (PCB) active at a time
 - Current state of process held in PCB:
 - “snapshot” of the execution and protection environment
 - Process needs CPU, resources
- Give out CPU time to different processes (Scheduling):
 - Only one process “running” at a time
 - Give more time to important processes
- Give pieces of resources to different processes (Protection):
 - Controlled access to non-CPU resources
 - E.g. Memory Mapping: Give each process their own address space

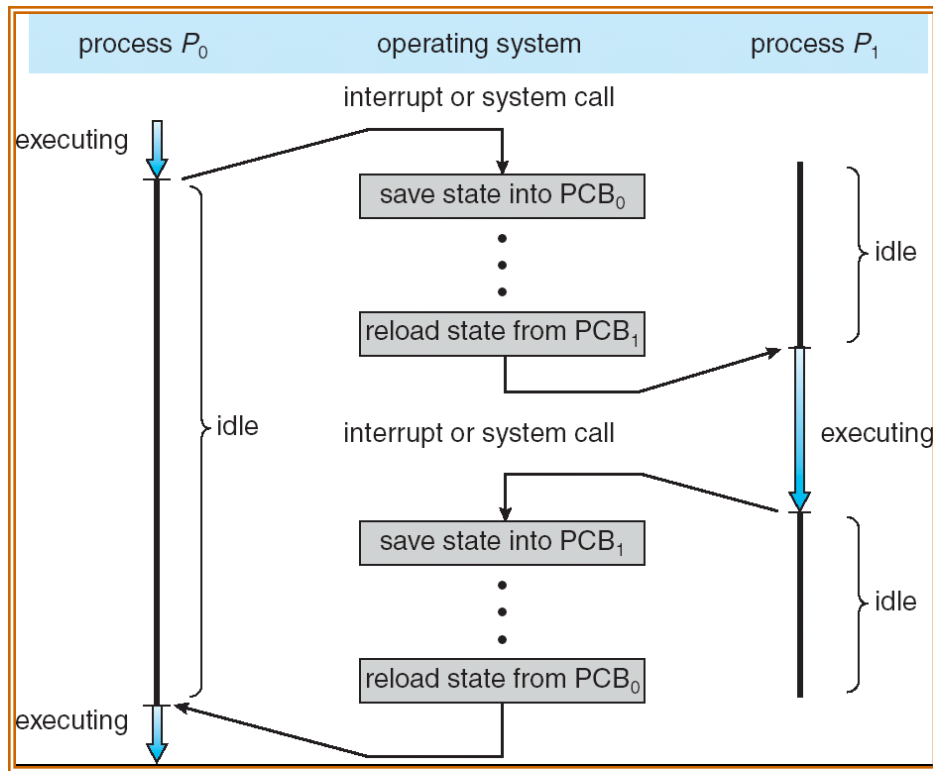


Process
Control
Block

Enabling Concurrency: Context Switch

- Task that switches CPU from one process to another process
 - the CPU must save the PCB state of the old process and load the saved PCB state of the new process.
- Context-switch time is overhead
 - System does no useful work while switching
 - Overhead sets minimum practical switching time; can become a bottleneck
- Time for context switch is dependent on hardware support (1- 1000 microseconds).

CPU Switch From Process to Process




- Code executed in kernel above is overhead
 - Overhead sets minimum practical switching time

Schedulers

- **Long-term scheduler (or job scheduler) -**
 - ❑ selects which processes should be brought into the ready queue.
 - ❑ invoked very infrequently (seconds, minutes); may be slow.
 - ❑ controls the degree of multiprogramming
- **Short term scheduler (or CPU scheduler) -**
 - ❑ selects which process should execute next and allocates CPU.
 - ❑ invoked very frequently (milliseconds) - must be very fast
- **Medium Term Scheduler**
 - ❑ swaps out process temporarily
 - ❑ balances load for better throughput

Medium Term (Time-sharing) Scheduler



Process Profiles

■ I/O bound process -

- spends more time in I/O, short CPU bursts, CPU underutilized.

■ CPU bound process -

- spends more time doing computations; few very long CPU bursts, I/O underutilized.


■ The right job mix:

- Long term scheduler - admits jobs to keep load balanced between I/O and CPU bound processes
- Medium term scheduler – ensures the right mix (by sometimes swapping out jobs and resuming them later)

Process Creation

- Processes are created and deleted dynamically
 - Process which creates another process is called a *parent* process; the created process is called a *child* process.
 - Result is a tree of processes
 - e.g. UNIX - processes have dependencies and form a hierarchy.
 - Resources required when creating process
 - CPU time, files, memory, I/O devices etc.
-

UNIX Process Hierarchy



•

What does it take to create a process?

- Must construct new PCB
 - Inexpensive
- Must set up new page tables for address space
 - More expensive
- Copy data from parent process? (Unix `fork()`)
 - Semantics of Unix `fork()` are that the child process gets a complete copy of the parent memory and I/O state
 - Originally very expensive
 - Much less expensive with “copy on write”
- Copy I/O state (file handles, etc)
 - Medium expense

Process Creation

■ Resource sharing

- ❑ Parent and children share all resources.
- ❑ Children share subset of parent's resources - prevents many processes from overloading the system.
- ❑ Parent and children share no resources.

■ Execution

- ❑ Parent and child execute concurrently.
- ❑ Parent waits until child has terminated.

■ Address Space

- ❑ Child process is duplicate of parent process.
- ❑ Child process has a program loaded into it.

UNIX Process Creation

- Fork system call creates new processes
 - `execve` system call is used after a fork to replace the processes memory space with a new program.
-

Process Termination

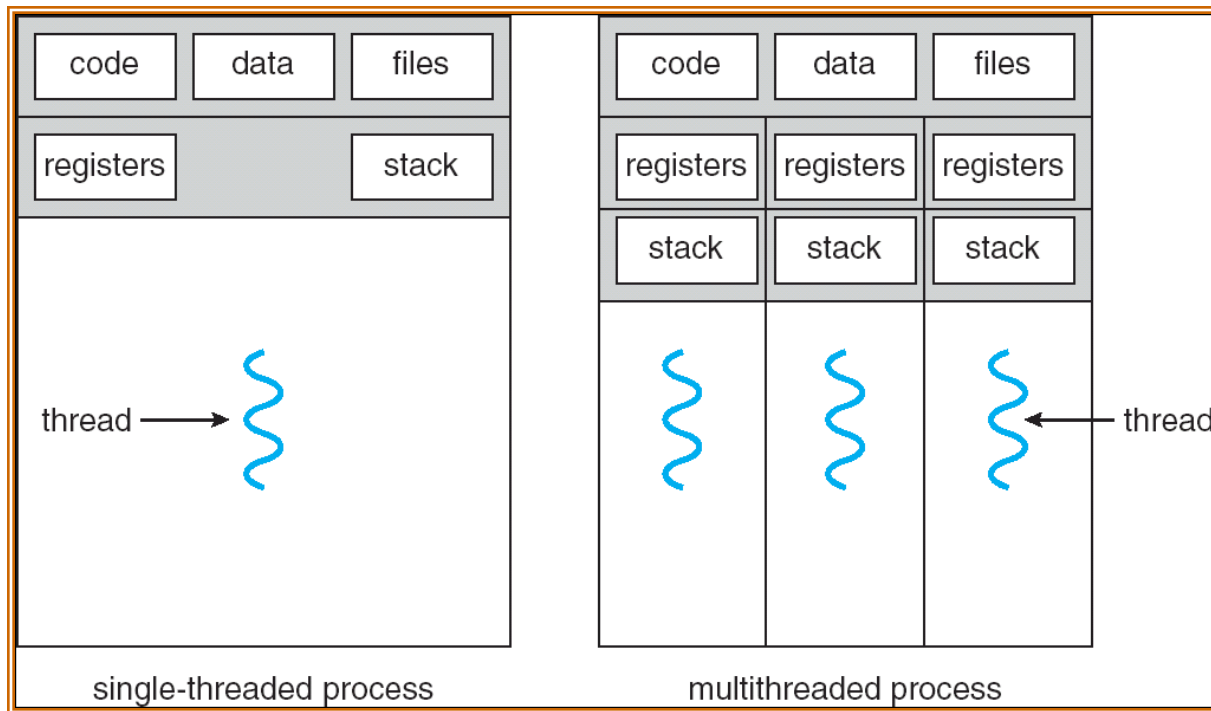
- Process executes last statement and asks the operating system to delete it (*exit*).
 - Output data from child to parent (via wait).
 - Process' resources are deallocated by operating system.
- Parent may terminate execution of child processes.
 - Child has exceeded allocated resources.
 - Task assigned to child is no longer required.
 - Parent is exiting
 - OS does not allow child to continue if parent terminates
 - Cascading termination

Threads

- Processes do not share resources well
 - high context switching overhead
- Idea: Separate concurrency from protection
- **Multithreading:** *a single program made up of a number of different concurrent activities*
- A thread (or lightweight process)
 - basic unit of CPU utilization; it consists of:
 - program counter, register set and stack space
 - A thread shares the following with peer threads:
 - code section, data section and OS resources (open files, signals)
 - No protection between threads
 - Collectively called a task.
- Heavyweight process is a task with one thread.



Single and Multithreaded Processes



- Threads encapsulate concurrency: “Active” component
- Address spaces encapsulate protection: “Passive” part
 - Keeps buggy program from trashing the system

Benefits

- Responsiveness
 - Resource Sharing
 - Economy
 - Utilization of MP Architectures
-

Threads(Cont.)

- In a multiple threaded task, while one server thread is blocked and waiting, a second thread in the same task can run.
 - Cooperation of multiple threads in the same job confers higher throughput and improved performance.
 - Applications that require sharing a common buffer (i.e. producer-consumer) benefit from thread utilization.
- Threads provide a mechanism that allows sequential processes to make blocking system calls while also achieving parallelism.

Thread State

- State shared by all threads in process/addr space
 - Contents of memory (global variables, heap)
 - I/O state (file system, network connections, etc)
 - State “private” to each thread
 - Kept in TCB \equiv Thread Control Block
 - CPU registers (including, program counter)
 - Execution stack
 - Parameters, Temporary variables
 - return PCs are kept while called procedures are executing
-

Threads (cont.)

- Thread context switch still requires a register set switch, but no memory management related work!!
 - Thread states -
 - *ready, blocked, running, terminated*
 - Threads share CPU and only one thread can run at a time.
 - No protection among threads.
-

Examples: Multithreaded programs

■ Embedded systems

- ❑ Elevators, Planes, Medical systems, Wristwatches
- ❑ Single Program, concurrent operations

■ Most modern OS kernels

- ❑ Internally concurrent because have to deal with concurrent requests by multiple users
- ❑ But no protection needed within kernel

■ Database Servers

- ❑ Access to shared data by many concurrent users
 - ❑ Also background utility processing must be done
-

More Examples: Multithreaded programs

■ Network Servers

- ❑ Concurrent requests from network
- ❑ Again, single program, multiple concurrent operations
- ❑ File server, Web server, and airline reservation systems

■ Parallel Programming (More than one physical CPU)

- ❑ Split program into multiple threads for parallelism
 - ❑ This is called Multiprocessing
-

# threads Per AS: # of addr spaces:	One	Many
One	MS/DOS, early Macintosh	Traditional UNIX
Many	Embedded systems (Geoworks, VxWorks, JavaOS,etc) JavaOS, Pilot(PC)	Mach, OS/2, Linux Windows 9x??? Win NT to XP, Solaris, HP-UX, OS X

Real operating systems have either

- ❑ One or many address spaces
- ❑ One or many threads per address space

Types of Threads

- Kernel-supported threads (Mach and OS/2)
 - User-level threads
 - Hybrid approach implements both user-level and kernel-supported threads (Solaris 2).
-

Kernel Threads

■ Supported by the Kernel

- Native threads supported directly by the kernel
- Every thread can run or block independently
- One process may have several threads waiting on different things

■ Downside of kernel threads: a bit expensive

- Need to make a crossing into kernel mode to schedule

■ Examples

- Windows XP/2000, Solaris, Linux, Tru64 UNIX, Mac OS X, Mach, OS/2
-

User Threads


- Supported above the kernel, via a set of library calls at the user level.
 - Thread management done by user-level threads library
 - User program provides scheduler and thread package
 - May have several user threads per kernel thread
 - User threads may be scheduled non-preemptively relative to each other (only switch on yield())
 - Advantages
 - Cheap, Fast
 - Threads do not need to call OS and cause interrupts to kernel
 - Disadv: If kernel is single threaded, system call from any thread can block the entire task.
- Example thread libraries:
 - POSIX Pthreads, Win32 threads, Java threads

Multithreading Models

- Many-to-One
 - One-to-One
 - Many-to-Many
-


Many-to-One

- Many user-level threads mapped to single kernel thread
- Examples:
 - ❑ Solaris Green Threads
 - ❑ GNU Portable Threads



One-to-One

- Each user-level thread maps to kernel thread




Examples

- Windows NT/XP/2000; Linux; Solaris 9 and later

Many-to-Many Model

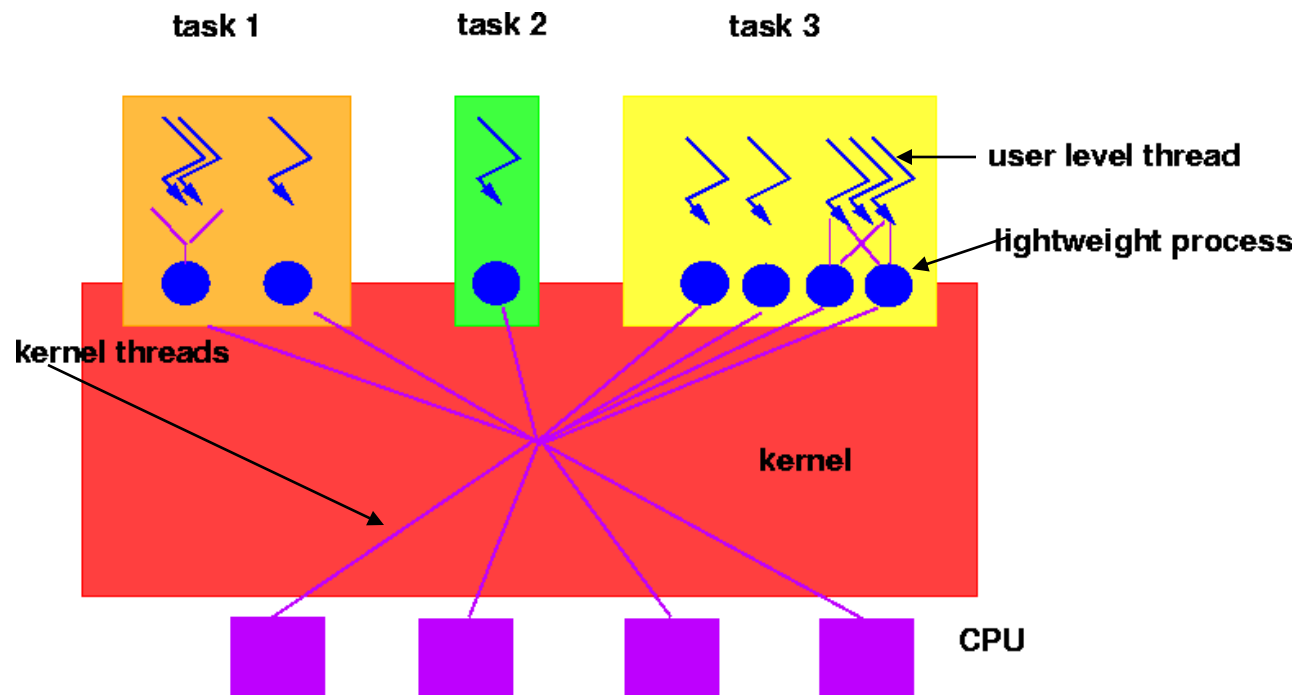
- Allows many user level threads to be mapped to many kernel threads
- Allows the operating system to create a sufficient number of kernel threads
- Solaris prior to version 9
- Windows NT/2000 with the *ThreadFiber* package



Thread Support in Solaris 2

- Solaris 2 is a version of UNIX with support for
 - kernel and user level threads, symmetric multiprocessing and real-time scheduling.
- Lightweight Processes (LWP)
 - intermediate between user and kernel level threads
 - each LWP is connected to exactly one kernel thread

Threads in Solaris 2




Threads in Solaris 2

- **Resource requirements of thread types**
 - Kernel Thread: small data structure and stack; thread switching does not require changing memory access information - relatively fast.
 - Lightweight Process: PCB with register data, accounting and memory information - switching between LWP is relatively slow.
 - User-level thread: only needs stack and program counter; no kernel involvement means fast switching. Kernel only sees the LWPs that support user-level threads.
-

Two-level Model

- Similar to M:M, except that it allows a user thread to be **bound** to kernel thread
- Examples
 - ▣ IRIX, HP-UX, Tru64 UNIX, Solaris 8 and earlier



Threading Issues

- Semantics of **fork()** and **exec()** system calls
 - Thread cancellation
 - Signal handling
 - Thread pools
 - Thread specific data
 - Scheduler activations
-


Multi(processing, programming, threading)

■ Definitions:

- ❑ Multiprocessing \equiv Multiple CPUs
- ❑ Multiprogramming \equiv Multiple Jobs or Processes
- ❑ Multithreading \equiv Multiple threads per Process

■ What does it mean to run two threads “concurrently”?

- ❑ Scheduler is free to run threads in any order and interleaving: FIFO, Random, ...
- ❑ Dispatcher can choose to run each thread to completion or time-slice in big chunks or small chunks



Cooperating Processes

■ Concurrent Processes can be

- Independent processes
 - cannot affect or be affected by the execution of another process.
- Cooperating processes
 - can affect or be affected by the execution of another process.


■ Advantages of process cooperation:

- Information sharing
- Computation speedup
- Modularity
- Convenience(e.g. editing, printing, compiling)

■ Concurrent execution requires


- process communication and process synchronization

Interprocess Communication (IPC)



- Separate address space isolates processes
 - High Creation/Memory Overhead; (Relatively) High Context-Switch Overhead
- Mechanism for processes to communicate and synchronize actions.
 - Via shared memory - Accomplished by mapping addresses to common DRAM
 - Read and Write through memory
 - Via Messaging system - processes communicate without resorting to shared variables.
 - `send()` and `receive()` messages
 - Can be used over the network!
 - Messaging system and shared memory not mutually exclusive
 - can be used simultaneously within a single OS or a single process.

Shared Memory Communication



- Communication occurs by “simply” reading/writing to shared address page
 - Really low overhead communication
 - Introduces complex synchronization problems

Cooperating Processes via Message Passing

- IPC facility provides two operations.

send(*message*) - message size can be fixed or variable

receive(*message*)

- If processes P and Q wish to communicate, they need to:

- ❑ establish a communication link between them
- ❑ exchange messages via send/receive

- Fixed vs. Variable size message

- ❑ Fixed message size - straightforward physical implementation, programming task is difficult due to fragmentation
- ❑ Variable message size - simpler programming, more complex physical implementation.

Implementation Questions

- ❑ How are links established?
- ❑ Can a link be associated with more than 2 processes?
- ❑ How many links can there be between every pair of communicating processes?
- ❑ What is the capacity of a link?
- ❑ Fixed or variable size messages?
- ❑ Unidirectional or bidirectional links?

.....

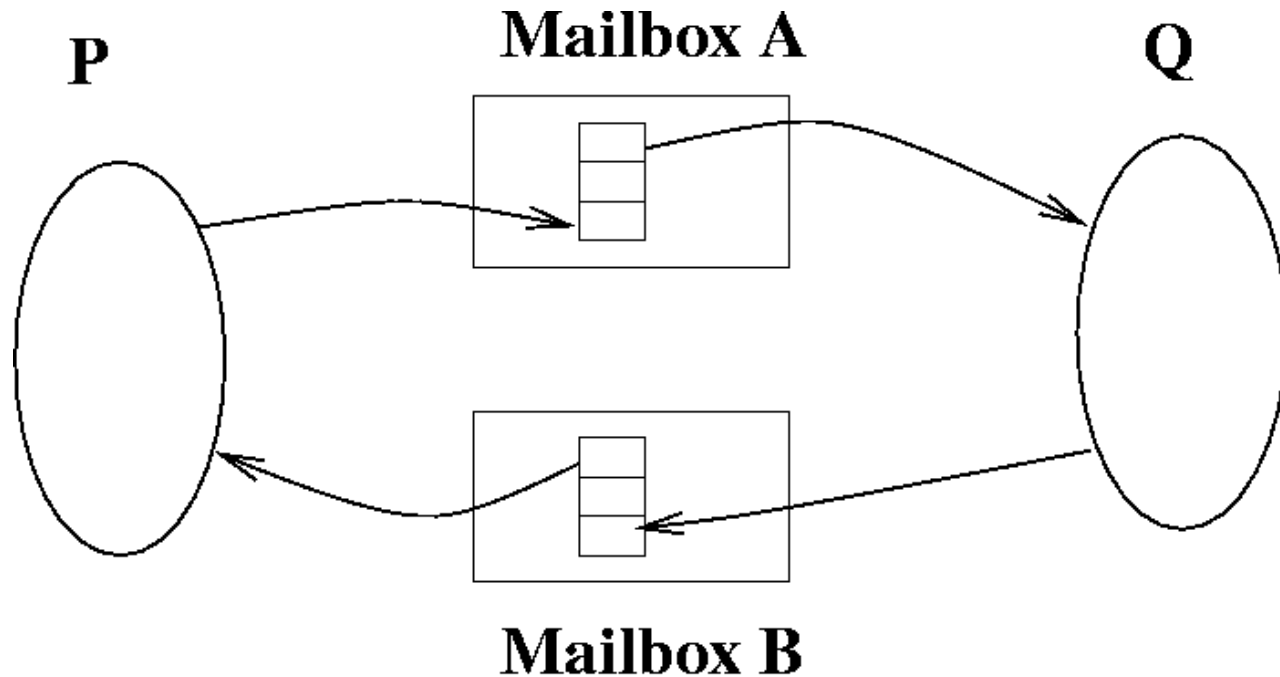
Direct Communication

- Sender and Receiver processes must name each other explicitly:
 - ❑ **send**(P , *message*) - send a message to process P
 - ❑ **receive**(Q , *message*) - receive a message from process Q
- Properties of communication link:
 - ❑ Links are established automatically.
 - ❑ A link is associated with exactly one pair of communicating processes.
 - ❑ Exactly one link between each pair.
 - ❑ Link may be unidirectional, usually bidirectional.

Indirect Communication

- Messages are directed to and received from mailboxes (also called ports)
 - Unique ID for every mailbox.
 - Processes can communicate only if they share a mailbox.
Send(*A, message*) /* send *message* to mailbox *A* */
Receive(*A, message*) /* receive *message* from mailbox *A* */
- Properties of communication link
 - Link established only if processes share a common mailbox.
 - Link can be associated with many processes.
 - Pair of processes may share several communication links
 - Links may be unidirectional or bidirectional

Indirect Communication using mailboxes



•

Mailboxes (cont.)

■ Operations

- create a new mailbox
- send/receive messages through mailbox
- destroy a mailbox

■ Issue: Mailbox sharing

- P1, P2 and P3 share mailbox A.
- P1 sends message, P2 and P3 receive... who gets message??

■ Possible Solutions

- disallow links between more than 2 processes
- allow only one process at a time to execute receive operation
- allow system to arbitrarily select receiver and then notify sender.

Message Buffering

- Link has some capacity - determine the number of messages that can reside temporarily in it.
- Queue of messages attached to link
 - Zero-capacity Queues: 0 messages
 - sender waits for receiver (synchronization is called *rendezvous*)
 - Bounded capacity Queues: Finite length of n messages
 - sender waits if link is full
 - Unbounded capacity Queues: Infinite queue length
 - sender never waits

Message Problems - Exception Conditions

❑ Process Termination

- Problem: P(sender) terminates, Q(receiver) blocks forever.

❑ Solutions:

- System terminates Q.
- System notifies Q that P has terminated.
- Q has an internal mechanism(timer) that determines how long to wait for a message from P.

- Problem: P(sender) sends message, Q(receiver) terminates. In automatic buffering, P sends message until buffer is full or forever. In no-buffering scheme, P blocks forever.

❑ Solutions:

- System notifies P
- System terminates P
- P and Q use acknowledgement with timeout

Message Problems - Exception Conditions

■ Lost Messages

- OS guarantees retransmission
- sender is responsible for detecting it using timeouts
- sender gets an exception

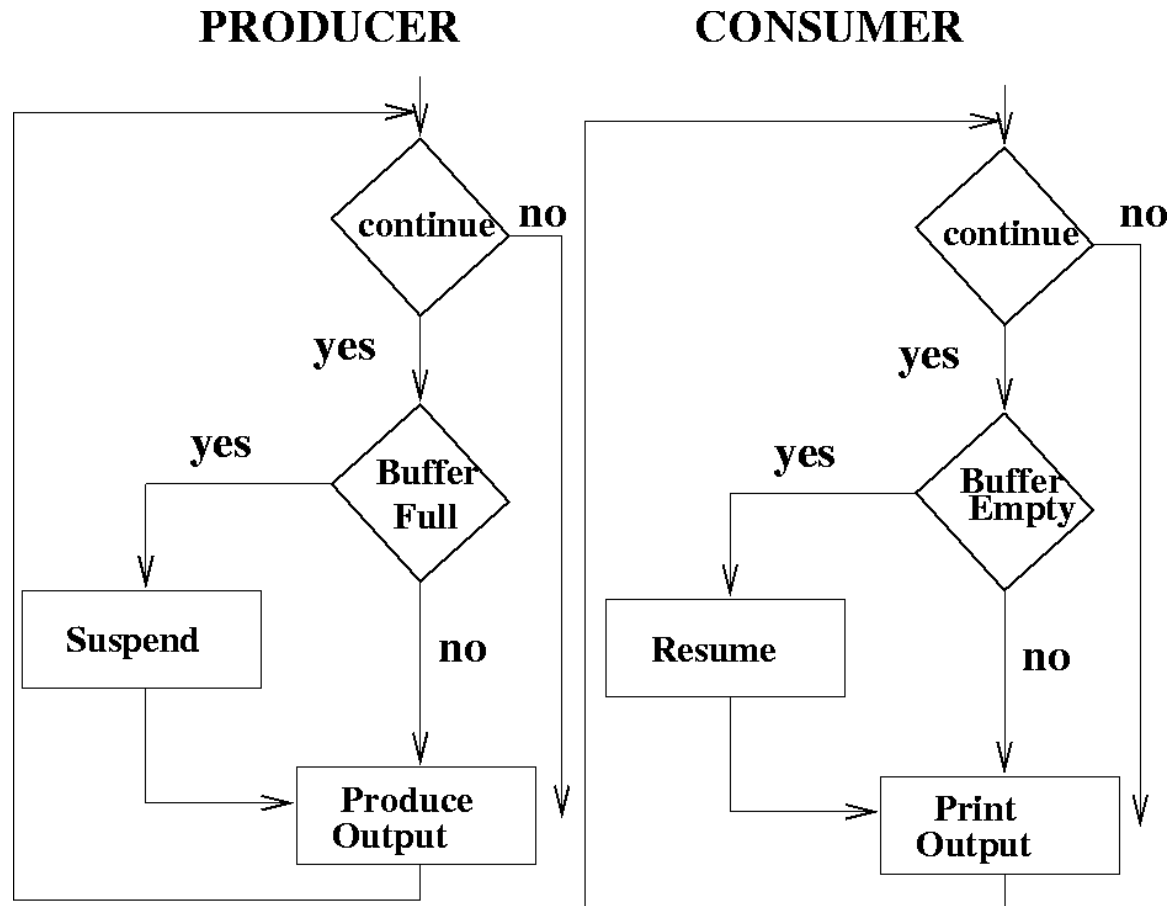
■ Scrambled Messages

- Message arrives from sender P to receiver Q, but information in message is corrupted due to noise in communication channel.
- Solution
 - need error detection mechanism, e.g. CHECKSUM
 - need error correction mechanism, e.g. retransmission

Producer-Consumer Problem

- Paradigm for cooperating processes;
 - producer process produces information that is consumed by a consumer process.
- We need buffer of items that can be filled by producer and emptied by consumer.
 - Unbounded-buffer places no practical limit on the size of the buffer. Consumer may wait, producer never waits.
 - Bounded-buffer assumes that there is a fixed buffer size. Consumer waits for new item, producer waits if buffer is full.
- Producer and Consumer must synchronize.

Producer-Consumer Problem



Bounded Buffer using IPC (messaging)

- Producer

repeat

...

produce an item in *nextp*;

...

send(*consumer*, *nextp*);

until false;

- Consumer

repeat

receive(*producer*, *nextc*);

...

consume item from *nextc*;

...

until false;

Bounded-buffer - Shared Memory Solution

■ Shared data

```
var n;  
type item = ....;  
var buffer: array[0..n-1] of item;  
in, out: 0..n-1;  
in := 0; out := 0; /* shared buffer = circular array */  
/* Buffer empty if in == out */  
/* Buffer full if (in+1) mod n == out */  
/* noop means 'do nothing' */
```

Bounded Buffer - Shared Memory Solution

- Producer process - creates filled buffers

repeat

...

produce an item in *nextp*

...

while $in+1 \bmod n = out$ **do** *noop*;

buffer[in] := *nextp*;

in := $in+1 \bmod n$;

until *false*;

Bounded Buffer - Shared Memory Solution

■ Consumer process - Empties filled buffers

repeat

while $in = out$ **do** *noop*;

$nextc := buffer[out]$;

$out := out + 1 \bmod n$;

...

consume the next item in *nextc*

...

until *false*