

ICS 143 - Principles of Operating Systems

Lecture 6 and 7 - Process Synchronization

Prof. Nalini Venkatasubramanian

nalini@ics.uci.edu

Outline

- Cooperating Processes
 - The Bounded Buffer Producer-Consumer Problem
 - The Critical Section Problem
 - Synchronization Hardware
 - Semaphores
 - Classical Problems of Synchronization
 - Critical Regions
 - Monitors
-


Cooperating Processes

- **Concurrent Processes can be**
 - Independent processes
 - cannot affect or be affected by the execution of another process.
 - Cooperating processes
 - can affect or be affected by the execution of another process.
 - **Advantages of process cooperation:**
 - Information sharing
 - Computation speedup
 - Modularity
 - Convenience(e.g. editing, printing, compiling)
 - **Concurrent execution requires**
 - process communication and process synchronization
-

Producer-Consumer Problem

- Paradigm for cooperating processes;
 - producer process produces information that is consumed by a consumer process.
 - We need buffer of items that can be filled by producer and emptied by consumer.
 - Unbounded-buffer places no practical limit on the size of the buffer. Consumer may wait, producer never waits.
 - Bounded-buffer assumes that there is a fixed buffer size. Consumer waits for new item, producer waits if buffer is full.
 - Producer and Consumer must synchronize.
-

Producer-Consumer Problem



Bounded-buffer - Shared Memory Solution

■ Shared data

```
var n;  
type item = ....;  
var buffer: array[0..n-1] of item;  
in, out: 0..n-1;  
in := 0; out := 0; /* shared buffer = circular array */  
/* Buffer empty if in == out */  
/* Buffer full if (in+1) mod n == out */  
/* noop means 'do nothing' */
```

Bounded Buffer - Shared Memory Solution

- Producer process - creates filled buffers

repeat

...

produce an item in *nextp*

...

while $in+1 \bmod n = out$ **do** *noop*;

buffer[in] := nextp;

in := in+1 mod n;

until *false*;

Bounded Buffer - Shared Memory Solution

■ Consumer process - Empties filled buffers

repeat

while $in = out$ **do** *noop*;

$nextc := buffer[out]$;

$out := out + 1 \bmod n$;

 ...

 consume the next item in *nextc*

 ...

until *false*

Background

- Concurrent access to shared data may result in data inconsistency.
- Maintaining data consistency requires mechanisms to ensure the orderly execution of cooperating processes.
- Shared memory solution to the bounded-buffer problem allows at most $(n-1)$ items in the buffer at the same time.

Bounded Buffer

- A solution that uses all N buffers is not that simple.

- Modify producer-consumer code by adding a variable *counter*, initialized to 0, incremented each time a new item is added to the buffer

- Shared data

```
type item = ....;  
var buffer: array[0..n-1] of item;  
in, out: 0..n-1;  
counter: 0..n;  
in, out, counter := 0;
```

Bounded Buffer

■ Producer process - creates filled buffers

repeat

...

produce an item in *nextp*

...

while *counter* = *n* **do** *noop*;

buffer[in] := *nextp*;

in := *in* + 1 **mod** *n*;

counter := *counter* + 1;

until *false*;

Bounded Buffer

■ Consumer process - Empties filled buffers

repeat

while *counter* = 0 **do** *noop*;

nextc := *buffer*[*out*];

out := *out* + 1 **mod** *n*;

counter := *counter* - 1;

 ...

 consume the next item in *nextc*

 ...

until *false*;

■ The statements

counter := *counter* + 1;

counter := *counter* - 1;

 must be executed *atomically*.

The Critical-Section Problem

- N processes all competing to use shared data.

- Structure of process P_i --- Each process has a code segment, called the critical section, in which the shared data is accessed.

repeat

entry section /* enter critical section */

critical section /* access shared variables */

exit section /* leave critical section */

remainder section /* do other work */

until false

- Problem

- Ensure that when one process is executing in its critical section, no other process is allowed to execute in its critical section.

Solution: Critical Section Problem - Requirements

- ❑ **Mutual Exclusion**

- ❑ If process P_i is executing in its critical section, then no other processes can be executing in their critical sections.

- ❑ **Progress**

- ❑ If no process is executing in its critical section and there exists some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely.

- ❑ **Bounded Waiting**

- ❑ A bound must exist on the number of times that other processes are allowed to enter their critical sections after a process has made a request to enter its critical section and before that request is granted.

Solution: Critical Section Problem - Requirements

- Assume that each process executes at a nonzero speed.
 - No assumption concerning relative speed of the n processes.
-

Solution: Critical Section Problem -- Initial Attempt

- Only 2 processes, P0 and P1
- General structure of process P_i (P_j)

repeat

entry section

critical section

exit section

remainder section

until false

- Processes may share some common variables to synchronize their actions.
-

Algorithm 1

- Shared Variables:

- **var** *turn*: (0..1);
initially *turn* = 0;
- *turn* = *i* \checkmark P_i can enter its critical section

- Process P_i

repeat

while *turn* \neq *i* do no-op;

critical section

turn := *j*;

remainder section

until false

Satisfies mutual exclusion, but not progress.

Algorithm 2

□ Shared Variables

- **var** *flag*: **array** (0..1) **of** boolean;
initially *flag*[0] = *flag*[1] = false;
- *flag*[*i*] = true ✓ P_i ready to enter its critical section

□ Process P_i

repeat

flag[*i*] := true;

while *flag*[*j*] **do** no-op;

critical section

flag[*i*] := false;

remainder section

until false

Can block indefinitely.... Progress requirement not met.

Algorithm 3

❑ Shared Variables

- **var** *flag*: **array** (0..1) **of** boolean;
initially *flag*[0] = *flag*[1] = false;
- *flag*[*i*] = true ☑ P_i ready to enter its critical section

❑ Process P_i

repeat

while *flag*[*j*] **do** no-op;
 flag[*i*] := true;

 critical section

flag[*i*] := false;

 remainder section

until false

Does not satisfy mutual exclusion requirement

Algorithm 4

- ❑ Combined Shared Variables of algorithms 1 and 2
- ❑ Process P_i

repeat

flag[i] := true;

turn := j;

while (*flag[j] and turn=j*) **do** *no-op;*

critical section

flag[i] := false;

remainder section

until false

YES!!! Meets all three requirements, solves the critical section problem for 2 processes.

Bakery Algorithm

■ Critical section for n processes

- Before entering its critical section, process receives a number. Holder of the smallest number enters critical section.
- If processes P_i and P_j receive the same number,
 - if $i \leq j$, then P_i is served first; else P_j is served first.
- The numbering scheme always generates numbers in increasing order of enumeration; i.e. 1,2,3,3,3,3,4,4,5,5

Bakery Algorithm (cont.)

■ Notation -

- Lexicographic order(ticket#, process id#)
 - $(a,b) < (c,d)$ if $(a < c)$ or if $((a = c) \text{ and } (b < d))$
 - $\max(a_0, \dots, a_{n-1})$ is a number, k , such that $k \geq a_i$ for $i = 0, \dots, \underline{n}-1$

■ Shared Data

var *choosing*: **array**[0..*n*-1] **of** *boolean*; (initialized to *false*)
number: **array**[0..*n*-1] **of** *integer*; (initialized to 0)

Bakery Algorithm (cont.)

repeat

choosing[*i*] := *true*;

number[*i*] := max(*number*[0], *number*[1], ..., *number*[*n*-1]) + 1;

choosing[*i*] := *false*;

for *j* := 0 **to** *n*-1

do begin

while *choosing*[*j*] **do** *no-op*;

while *number*[*j*] <> 0

and (*number*[*j*], *j*) < (*number*[*i*], *i*) **do** *no-op*;

end;

critical section

number[*i*] := 0;

remainder section

until *false*;

Hardware Solutions for Synchronization

- Mutual exclusion solutions presented depend on memory hardware having read/write cycle.
 - If multiple reads/writes could occur to the same memory location at the same time, this would not work.
 - Processors with caches but no cache coherency cannot use the solutions
 - In general, it is impossible to build mutual exclusion without a primitive that provides some form of mutual exclusion.
 - How can this be done in the hardware???
-

Synchronization Hardware

- Test and modify the content of a word atomically - **Test-and-set** instruction

```
function Test-and-Set (var target: boolean): boolean;  
  begin  
    Test-and-Set := target;  
    target := true;  
  end;
```

- Similarly **“SWAP”** instruction

Mutual Exclusion with Test-and-Set

- Shared data: var lock: boolean (initially false)

- Process P_i

repeat

while *Test-and-Set*(lock) do no-op;

critical section

lock := false;

remainder section

until false;

Bounded Waiting Mutual Exclusion with Test-and-Set

```
var  $j : 0..n-1$ ;  
     $key : boolean$ ;  
repeat  
     $waiting[i] := true$ ;  $key := true$ ;  
    while  $waiting[i]$  and  $key$  do  $key := Test\text{-}and\text{-}Set(lock)$ ;  
     $waiting[i] := false$ ;  
    critical section  
     $j := j + 1 \bmod n$ ;  
    while  $(j \neq i)$  and (not  $waiting[j]$ ) do  $j := j + 1 \bmod n$ ;  
    if  $j = i$  then  $lock := false$ ;  
        else  $waiting[j] := false$ ;  
    remainder section  
  
until  $false$ ;
```

Semaphore

- Semaphore S - integer variable

- used to represent number of abstract resources

- Can only be accessed via two indivisible (atomic) operations

wait (S): **while** $S \leq 0$ **do** no-op
 $S := S - 1$;

signal (S): $S := S + 1$;

- P or *wait* used to acquire a resource, decrements count
 - V or *signal* releases a resource and increments count
 - If P is performed on a *count* ≤ 0 , process must wait for V or the release of a resource.

Example: Critical Section for n Processes

- Shared variables

var mutex: semaphore
initially *mutex* = 1

- Process P_i

repeat

wait(mutex);

critical section

signal (mutex);

remainder section

until false

Semaphore as a General Synchronization Tool

- Execute B in P_j only after A execute in P_i
- Use semaphore *flag* initialized to 0
- Code:

P_i	P_j
⋮	⋮
A	$wait(flag)$
$signal(flag)$	B

Problem...

- Busy Waiting, uses CPU that others could use. This type of semaphore is called a *spinlock*.
 - OK for short times since it prevents a context switch.
 - For longer runtimes, need to modify P and V so that processes can *block* and *resume*.
-

Semaphore Implementation

- Define a semaphore as a record

```
type semaphore = record  
    value: integer;  
    L: list of processes;  
end;
```

- Assume two simple operations

- *block* suspends the process that invokes it.
- *wakeup(P)* resumes the execution of a blocked process *P*.

Semaphore Implementation(cont.)

- Semaphore operations are now defined as

wait (S): $S.value := S.value - 1;$

if $S.value < 0$

then begin

add this process to $S.L$;

block;

end;

signal (S): $S.value := S.value + 1;$

if $S.value \leq 0$

then begin

remove a process P from $S.L$;

wakeup(P);

end;

Block/Resume Semaphore Implementation

- If process is blocked, enqueue PCB of process and call scheduler to run a different process.
- Semaphores are executed atomically;
 - no two processes execute *wait* and *signal* at the same time.
 - Mutex can be used to make sure that two processes do not change count at the same time.
 - If an interrupt occurs while mutex is held, it will result in a long delay.
 - Solution: Turn off interrupts during critical section.

Deadlock and Starvation

- ❑ Deadlock - two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes.

- Let S and Q be semaphores initialized to 1

P_0
 $wait(S);$
 $wait(Q);$
 \vdots
 $signal(S);$
 $signal(Q);$

P_1
 $wait(Q);$
 $wait(S);$
 \vdots
 $signal(Q);$
 $signal(S);$

- ❑ Starvation- indefinite blocking. A process may never be removed from the semaphore queue in which it is suspended.

Two Types of Semaphores

- Counting Semaphore - integer value can range over an unrestricted domain.
 - Binary Semaphore - integer value can range only between 0 and 1; simpler to implement.
 - Can implement a counting semaphore S as a binary semaphore.
-

Implementing S (counting sem.) as a Binary Semaphore

■ Data Structures

```
var S1 : binary-semaphore;  
    S2 : binary-semaphore;  
    S3 : binary-semaphore;  
    C : integer;
```

■ Initialization

```
S1 = S3 = 1;  
S2 = 0;  
C = initial value of semaphore S;
```

Implementing S

Wait operation

```
wait(S3);  
wait(S1);  
C := C-1;  
if C < 0  
then begin  
    signal (S1);  
    wait(S2);  
end  
else signal (S1);  
    signal (S3);
```

Signal operation

```
wait(S1);  
C := C + 1;  
if C <= 0 then signal (S2);  
    signal (S1);
```

Classical Problems of Synchronization

- Bounded Buffer Problem
- Readers and Writers Problem
- Dining-Philosophers Problem

Bounded Buffer Problem

■ Shared data

```
type item = ....;  
var buffer: array[0..n-1] of item;  
full, empty, mutex : semaphore;  
nextp, nextc : item;  
full := 0; empty := n; mutex := 1;
```


Bounded Buffer Problem

■ Producer process - creates filled buffers

repeat

...

produce an item in *nextp*

...

wait (*empty*);

wait (*mutex*);

...

add *nextp* to buffer

...

signal (*mutex*);

signal (*full*);

until *false*;

Bounded Buffer Problem

■ Consumer process - Empties filled buffers

repeat

wait (full);

wait (mutex);

...

remove an item from *buffer* to *nextc*

...

signal (mutex);

signal (empty);

...

consume the next item in *nextc*

...

until *false*;

Readers-Writers Problem

■ Shared Data

```
var mutex, wrt: semaphore (=1);  
    readcount: integer (= 0);
```

■ Writer Process

```
wait(wrt);  
    ...  
    writing is performed  
    ...  
signal(wrt);
```

Readers-Writers Problem


■ Reader process

```
wait(mutex);
    readcount := readcount + 1;
    if readcount = 1 then wait(wrt);
    signal(mutex);

    ...
    reading is performed

    ...
wait(mutex);
    readcount := readcount - 1;
    if readcount = 0 then signal(wrt);
    signal(mutex);
```

Dining-Philosophers Problem



Shared Data

var *chopstick*: **array** [0..4] **of** *semaphore* (=1 initially);

Dining Philosophers Problem

□ Philosopher i :

repeat

wait (*chopstick*[i]);

wait (*chopstick*[$i+1 \bmod 5$]);

...

eat

...

signal (*chopstick*[i]);

signal (*chopstick*[$i+1 \bmod 5$]);

...

think

...

until *false*;

Higher Level Synchronization

- Timing errors are still possible with semaphores

- Example 1

- signal(mutex);*
 ...
 critical region

- wait(mutex);*

- Example 2

- wait(mutex);*
 ...
 critical region

- wait(mutex);*

- Example 3

- wait(mutex);*
 ...
 critical region

- Forgot to signal

Conditional Critical Regions

- High-level synchronization construct
- A shared variable v of type T is declared as:
var v : shared T
- Variable v is accessed only inside statement
region v when B do \underline{S}
where B is a boolean expression.
While statement S is being executed, no other process can access variable v .

Critical Regions (cont.)

- Regions referring to the same shared variable exclude each other in time.
- When a process tries to execute the region statement, the Boolean expression B is evaluated.
 - If B is true, statement S is executed.
 - If it is false, the process is delayed until B becomes true and no other process is in the region associated with v .

Example - Bounded Buffer

❑ Shared variables

var *buffer*. **shared record**

pool: **array**[0..*n*-1] **of** *item*;

count, in, out: *integer*;

end;

❑ Producer Process inserts *nextp* into the shared buffer

region *buffer* **when** *count* < *n*

do begin

pool[*in*] := *nextp*;

in := *in*+1 **mod** *n*;

count := *count* + 1;

end;

Bounded Buffer Example

- Consumer Process removes an item from the shared buffer and puts it in *nextc*

region *buffer* **when** *count* > 0

do begin

nextc := *pool*[*out*];

out := *out*+1 **mod** *n*;

count := *count* -1;

end;

Implementing Regions

- Region x when B do S

*var mutex, first-delay, second-delay: semaphore;
first-count, second-count: integer;*

- Mutually exclusive access to the critical section is provided by mutex.

If a process cannot enter the critical section because the Boolean expression B is false,

it initially waits on the first-delay semaphore;

moved to the second-delay semaphore before it is allowed to reevaluate B .

Implementation

- Keep track of the number of processes waiting on *first-delay* and *second-delay*, with *first-count* and *second-count* respectively.
- The algorithm assumes a FIFO ordering in the queueing of processes for a semaphore.
- For an arbitrary queueing discipline, a more complicated implementation is required.

Implementing Regions

```
wait(mutex);
while not B
do begin  first-count := first-count + 1;
         if second-count > 0
           then signal (second-delay);
           else signal (mutex);
         wait(first-delay);
         first-count := first-count - 1;
         second-count := second-count + 1;
         if first-count > 0 then signal (first-delay)
           else signal (second-delay);
         wait(second-delay);
         second-count := second-count - 1;
       end;
S;
if first-count > 0 then signal (first-delay);
  else if second-count > 0
    then signal (second-delay);
    else signal (mutex);
```

Monitors

High-level synchronization construct that allows the safe sharing of an abstract data type among concurrent processes.

```
type monitor-name = monitor
  variable declarations
  procedure entry P1 (...);
    begin ... end;
  procedure entry P2 (...);
    begin ... end;
    ⋮
  procedure entry Pn(...);
    begin ... end;
begin
  initialization code
end.
```

Monitors

- To allow a process to wait within the monitor, a condition variable must be declared, as:

var *x,y: condition*

- Condition variable can only be used within the operations *wait* and *signal*. Queue is associated with condition variable.

- The operation

x.wait;

means that the process invoking this operation is suspended until another process invokes

x.signal;

- The *x.signal* operation resumes exactly one suspended process. If no process is suspended, then the signal operation has no effect.

Dining Philosophers

```
type dining-philosophers= monitor
  var state: array[0..4] of (thinking, hungry, eating);
  var self: array[0..4] of condition;
  // condition where philosopher I can delay himself when hungry but    //
  // is unable to obtain chopstick(s)
  procedure entry pickup (i:0..4);
  begin
    state[i] := hungry;
    test(i); //test that your left and right neighbors are not eating
    if state [i] <> eating then self [i].wait;
  end;
  procedure entry putdown (i:0..4);
  begin
    state[i] := thinking;
    test (i + 4 mod 5 ); // signal left neighbor
    test (i + 1 mod 5 ); // signal right neighbor
  end;
```

Dining Philosophers (cont.)

```
procedure test (k :0..4);  
  begin  
    if state [k + 4 mod 5] <> eating  
      and state [k] = hungry  
      and state [k + 1 mod 5] <> eating  
    then  
      begin  
        state[k] := eating;  
        self [k].signal;  
      end;  
    end;  
  
  begin  
    for i := 0 to 4  
      do state[i] := thinking;  
    end;
```