

---

# ICS 143 - Principles of Operating Systems

---

Lectures 8 and 9 - Deadlocks

Prof. Nalini Venkatasubramanian

[nalini@ics.uci.edu](mailto:nalini@ics.uci.edu)

# Outline

- System Model
- Deadlock Characterization
- Methods for handling deadlocks
- Deadlock Prevention
- Deadlock Avoidance
- Deadlock Detection
- Recovery from Deadlock
- Combined Approach to Deadlock Handling

# The Deadlock Problem


- A set of blocked processes each holding a resource and waiting to acquire a resource held by another process in the set.
  - **Example 1**
    - System has 2 tape drives. P1 and P2 each hold one tape drive and each needs the other one.
  - **Example 2**
    - Semaphores A and B each initialized to 1

<i>P0</i>	<i>P1</i>
<i>wait(A)</i>	<i>wait(B)</i>
<i>wait(B)</i>	<i>wait(A)</i>

# Definitions

- A process is *deadlocked* if it is waiting for an event that will never occur.
  - Typically, more than one process will be involved in a deadlock (the deadly embrace).
- A process is *indefinitely postponed* if it is delayed repeatedly over a long period of time while the attention of the system is given to other processes,
  - i.e. the process is ready to proceed but never gets the CPU.

# Example - Bridge Crossing



- ❑ Assume traffic in one direction.
  - Each section of the bridge is viewed as a resource.
- ❑ If a deadlock occurs, it can be resolved only if one car backs up (preempt resources and rollback).
  - Several cars may have to be backed up if a deadlock occurs.
  - Starvation is possible

# Resources

## ■ Resource

- commodity required by a process to execute

## ■ Resources can be of several types

### ■ Serially Reusable Resources

- CPU cycles, memory space, I/O devices, files
- acquire -> use -> release

### ■ Consumable Resources

- Produced by a process, needed by a process - e.g. Messages, buffers of information, interrupts
- create -> acquire -> use
- Resource ceases to exist after it has been used

# System Model

- Resource types
  - $R_1, R_2, \dots, R_m$
- Each resource type  $R_i$  has  $W_i$  instances
- Assume serially reusable resources
  - request  $\rightarrow$  use  $\rightarrow$  release

# Conditions for Deadlock

- ❑ The following 4 conditions are necessary and sufficient for deadlock (must hold simultaneously)
  - Mutual Exclusion:
    - ❑ Only once process at a time can use the resource.
  - Hold and Wait:
    - ❑ Processes hold resources already allocated to them while waiting for other resources.
  - No preemption:
    - ❑ Resources are released by processes holding them only after that process has completed its task.
  - Circular wait:
    - ❑ A circular chain of processes exists in which each process waits for one or more resources held by the next process in the chain.



# Resource Allocation Graph


- A set of vertices  $V$  and a set of edges  $E$
- $V$  is partitioned into 2 types
  - $P = \{P_1, P_2, \dots, P_n\}$  - the set of processes in the system
  - $R = \{R_1, R_2, \dots, R_n\}$  - the set of resource types in the system
- Two kinds of edges
  - Request edge - Directed edge  $P_i \dashrightarrow R_j$
  - Assignment edge - Directed edge  $R_j \dashrightarrow P_i$

# Resource Allocation Graph


- Process




- Resource type with 4 instances




- $P_i$  requests instance of  $R_j$




- $P_i$  is holding an instance of  $R_j$




# Graph with no cycles







# Graph with cycles



# Graph with cycles and deadlock





# Basic facts

- If graph contains no cycles
  - NO DEADLOCK
- If graph contains a cycle
  - if only one instance per resource type, then deadlock
  - if several instances per resource type, possibility of deadlock.

# Methods for handling deadlocks

- Ensure that the system will never enter a deadlock state.
- Allow the system to potentially enter a deadlock state, detect it and then recover
- Ignore the problem and pretend that deadlocks never occur in the system;
  - Used by many operating systems, e.g. UNIX



# Deadlock Management

- ❑ Prevention

- ❑ Design the system in such a way that deadlocks can never occur

- ❑ Avoidance

- ❑ Impose less stringent conditions than for prevention, allowing the possibility of deadlock but sidestepping it as it occurs.

- ❑ Detection

- ❑ Allow possibility of deadlock, determine if deadlock has occurred and which processes and resources are involved.

- ❑ Recovery

- ❑ After detection, clear the problem, allow processes to complete and resources to be reused. May involve destroying and restarting processes.
-

# Deadlock Prevention

- ❑ If any one of the conditions for deadlock (with reusable resources) is denied, deadlock is impossible.
- ❑ Restrain ways in which requests can be made
  - Mutual Exclusion
    - ❑ non-issue for sharable resources
    - ❑ cannot deny this for non-sharable resources (important)
  - Hold and Wait - guarantee that when a process requests a resource, it does not hold other resources.
    - ❑ Force each process to acquire all the required resources at once. Process cannot proceed until all resources have been acquired.
    - ❑ Low resource utilization, starvation possible

# Deadlock Prevention (cont.)

## ■ No Preemption

- ❑ If a process that is holding some resources requests another resource that cannot be immediately allocated to it, the process releases the resources currently being held.
- ❑ Preempted resources are added to the list of resources for which the process is waiting.
- ❑ Process will be restarted only when it can regain its old resources as well as the new ones that it is requesting.

## ■ Circular Wait

- ❑ Impose a total ordering of all resource types.
- ❑ Require that processes request resources in increasing order of enumeration; if a resource of type  $N$  is held, process can only request resources of types  $> N$ .

---

# Deadlock Avoidance

- Set of resources, set of customers, banker
  - Rules
    - Each customer tells banker maximum number of resources it needs.
    - Customer borrows resources from banker.
    - Customer returns resources to banker.
    - Customer eventually pays back loan.
  - Banker only lends resources if the system will be in a *safe state* after the loan.
-

# Deadlock Avoidance

- Requires that the system has some additional apriori information available.
  - Simplest and most useful model requires that each process declare the maximum number of resources of each type that it may need.
  - The deadlock-avoidance algorithm dynamically examines the resource-allocation state to ensure that there can never be a circular-wait condition.
  - Resource allocation state is defined by the number of available and allocated resources, and the maximum demands of the processes.

# Safe state

- When a process requests an available resource, system must decide if immediate allocation leaves the system in a safe state.
- System is in safe state if there exists a safe sequence of all processes.
- Sequence  $\langle P_1, P_2, \dots, P_n \rangle$  is safe, if for each  $P_i$ , the resources that  $P_i$  can still request can be satisfied by currently available resources + resources held by  $P_j$  with  $j < i$ .
  - If  $P_i$  resource needs are not available,  $P_i$  can wait until all  $P_j$  have finished.
  - When  $P_j$  is finished,  $P_i$  can obtain needed resources, execute, return allocated resources, and terminate.
  - When  $P_i$  terminates,  $P_{i+1}$  can obtain its needed resources...

---

# Basic Facts

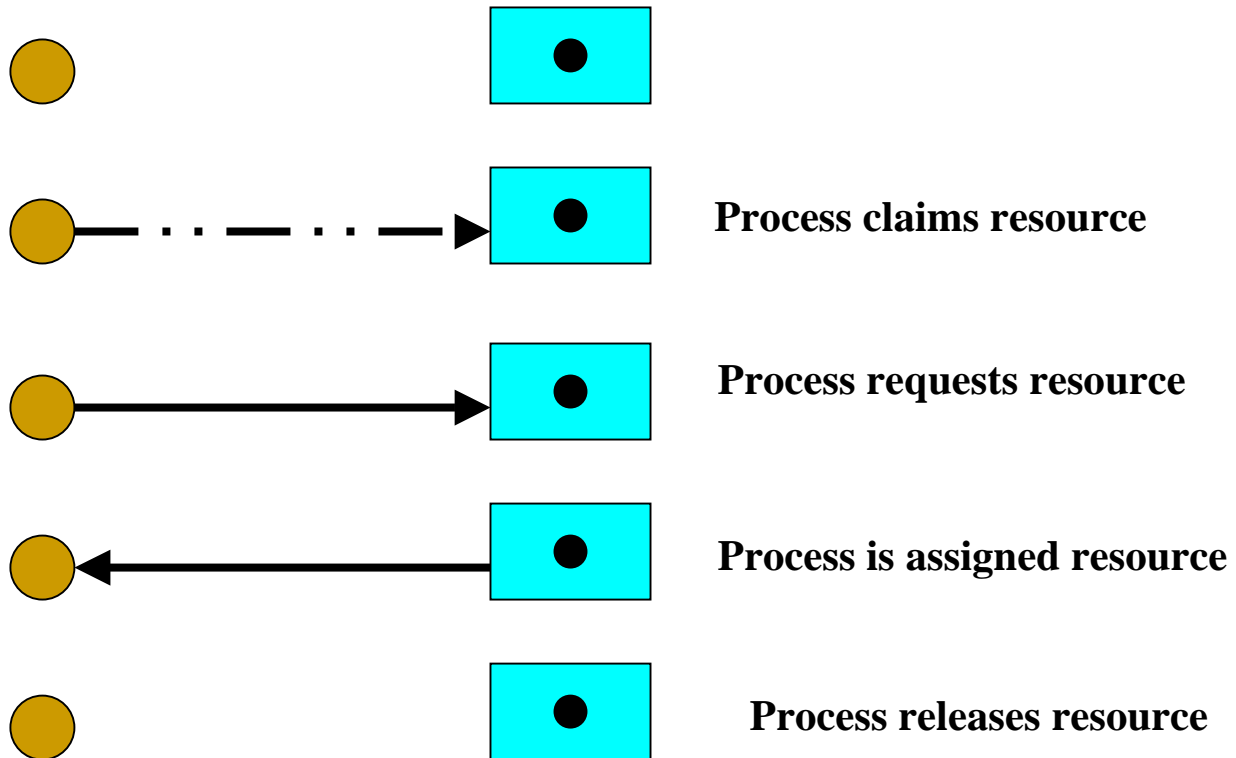
- If a system is in a safe state  $\Rightarrow$  no deadlocks.
  - If a system is in unsafe state  $\Rightarrow$  possibility of deadlock.
  - Avoidance  $\Rightarrow$  ensure that a system will never reach an unsafe state.
-

# Resource Allocation Graph Algorithm

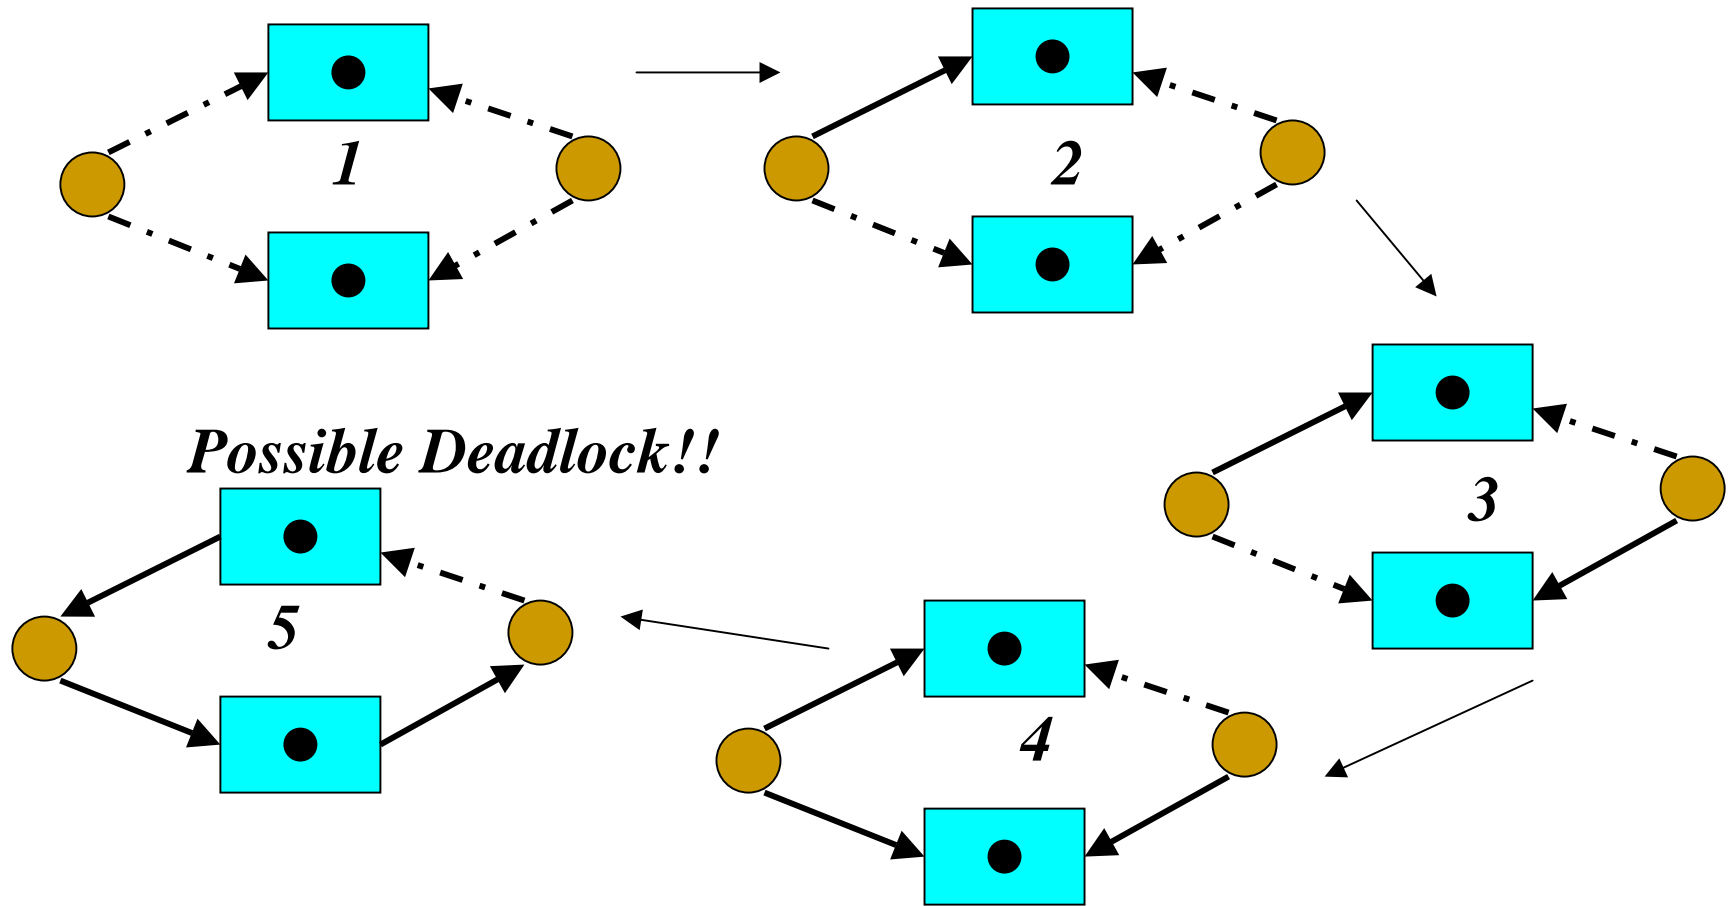
- Used for deadlock avoidance when there is only one instance of each resource type.
  - Claim edge:  $P_i \rightarrow R_j$  indicates that process  $P_i$  may request resource  $R_j$ ; represented by a dashed line.
  - Claim edge converts to request edge when a process requests a resource.
  - When a resource is released by a process, assignment edge reconverts to claim edge.
  - Resources must be claimed a priori in the system.
- If request assignment does not result in the formation of a cycle in the resource allocation graph - safe state, else unsafe state.



# Claim Graph



# Claim Graph



---

# Banker's Algorithm

- Used for multiple instances of each resource type.
  - Each process must a priori claim maximum use of each resource type.
  - When a process requests a resource it may have to wait.
  - When a process gets all its resources it must return them in a finite amount of time.
-

# Data Structures for the Banker's Algorithm

- Let  $n$  = number of processes and  $m$  = number of resource types.
  - *Available*: Vector of length  $m$ . If  $Available[j] = k$ , there are  $k$  instances of resource type  $R_j$  available.
  - *Max*:  $n \times m$  matrix. If  $Max[i,j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$ .
  - *Allocation*:  $n \times m$  matrix. If  $Allocation[i,j] = k$ , then process  $P_i$  is currently allocated  $k$  instances of resource type  $R_j$ .
  - *Need*:  $n \times m$  matrix. If  $Need[i,j] = k$ , then process  $P_i$  may need  $k$  more instances of resource type  $R_j$  to complete its task.

$$Need[i,j] = Max[i,j] - Allocation[i,j]$$

# Safety Algorithm

- Let *Work* and *Finish* be vectors of length *m* and *n*, respectively. Initialize
  - *Work* := *Available*
  - *Finish*[*i*] := *false* for *i* = 1, 2, ..., *n*.
- Find an *i* (i.e. process *P<sub>i</sub>*) such that both:
  - *Finish*[*i*] = *false*
  - *Need<sub>i</sub>* ≤ *Work*
  - If no such *i* exists, go to step 4.
- *Work* := *Work* + *Allocation<sub>i</sub>*
  - *Finish*[*i*] := *true*
  - go to step 2
- If *Finish*[*i*] = *true* for all *i*, then the system is in a safe state.

# Resource-Request Algorithm for Process $P_i$

- Request\_i = request vector for process  $P_i$ . If Request\_i[j] = k, then process  $P_i$  wants k instances of resource type  $R_j$ .
  - STEP 1: If  $Request(i) \leq Need(i)$ , go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim.
  - STEP 2: If  $Request(i) \leq Available$ , go to step 3. Otherwise,  $P_i$  must wait since resources are not available.
  - STEP 3: Pretend to allocate requested resources to  $P_i$  by modifying the state as follows:
    - $Available := Available - Request(i);$
    - $Allocation(i) := Allocation(i) + Request(i);$
    - $Need(i) := Need(i) - Request(i);$
  - If safe  $\Rightarrow$  resources are allocated to  $P_i$ .
  - If unsafe  $\Rightarrow P_i$  must wait and the old resource-allocation state is restored.

# Example of Banker's Algorithm

- 5 processes
  - P0 - P4;
- 3 resource types
  - A(10 instances), B (5 instances), C (7 instances)
- Snapshot at time T0

	Allocation			Max			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	5	3	3	3	2
P1	2	0	0	3	2	2			
P2	3	0	2	9	0	2			
P3	2	1	1	2	2	2			
P4	0	0	2	4	3	3			

## Example (cont.)

- The content of the matrix *Need* is defined to be *Max - Allocation*.
- The system is in a safe state since the sequence  $\langle P1, P3, P4, P2, P0 \rangle$  satisfies safety criteria.

	Need		
	A	B	C
P0	7	4	3
P1	1	2	2
P2	6	0	0
P3	0	1	1
P4	4	3	1



# Example: P1 requests (1,0,2)

- Check to see that Request  $\leq$  Available

- $((1,0,2) \leq (3,3,2)) \Rightarrow \text{true}.$

	Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	4	3	2	3	0
P1	3	0	2	0	2	0			
P2	3	0	2	6	0	0			
P3	2	1	1	0	1	1			
P4	0	0	2	4	3	1			

## Example (cont.)

- Executing the safety algorithm shows that sequence  $\langle P1, P3, P4, P0, P2 \rangle$  satisfies safety requirement.
- Can request for (3,3,0) by P4 be granted?
- Can request for (0,2,0) by P0 be granted?

---

# Deadlock Detection

- Allow system to enter deadlock state
  - Detection Algorithm
  - Recovery Scheme
-

# Single Instance of each resource type

- Maintain wait-for graph
    - Nodes are processes
    - $P_i \rightarrow P_j$  if  $P_i$  is waiting for  $P_j$ .
  - Periodically invoke an algorithm that searches for a cycle in the graph.
  - An algorithm to detect a cycle in a graph requires an order of  $n^2$  operations, where  $n$  is the number of vertices in the graph.
-

# Several instances of a resource type

## ■ Data Structures

- *Available*: Vector of length  $m$ . If  $Available[j] = k$ , there are  $k$  instances of resource type  $R_j$  available.
- *Allocation*:  $n \times m$  matrix. If  $Allocation[i,j] = k$ , then process  $P_i$  is currently allocated  $k$  instances of resource type  $R_j$ .
- *Request*: An  $n \times m$  matrix indicates the current request of each process. If  $Request[i,j] = k$ , then process  $P_i$  is requesting  $k$  more instances of resource type  $R_j$ .

# Deadlock Detection Algorithm

- Step 1: Let *Work* and *Finish* be vectors of length  $m$  and  $n$ , respectively. Initialize
  - $Work := Available$
  - For  $i = 1, 2, \dots, n$ , if  $Allocation(i) \neq 0$ , then  $Finish[i] := false$ , otherwise  $Finish[i] := true$ .
- Step 2: Find an index  $i$  such that both:
  - $Finish[i] = false$
  - $Request(i) \leq Work$
  - If no such  $i$  exists, go to step 4.

# Deadlock Detection Algorithm

- Step 3:  $Work := Work + Allocation(i)$ 
  - $Finish[i] := true$
  - go to step 2
- Step 4: If  $Finish[i] = false$  for some  $i$ ,  $1 \leq i \leq n$ , then the system is in a deadlock state. Moreover, if  $Finish[i] = false$ , then  $P_i$  is deadlocked.

Algorithm requires an order of  $m \times (n^2)$  operations to detect whether the system is in a deadlocked state.

# Example of Detection Algorithm

- 5 processes -  $P0 - P4$ ; 3 resource types -  $A(7$  instances),  $B(2$  instances),  $C(6$  instances)
- Snapshot at time  $T_0$ :  $\langle P0, P2, P3, P1, P4 \rangle$  will result in  $Finish[i] = \text{true}$  for all  $i$ .

	Allocation			Max			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	0	0	0	0	0	0
P1	2	0	0	2	0	2			
P2	3	0	3	0	0	0			
P3	2	1	1	1	0	0			
P4	0	0	2	0	0	2			



# Example (cont.)

- P2 requests an additional instance of type C.
- State of system
  - Can reclaim resources held by process P0, but insufficient resources to fulfill other processes' requests.
  - Deadlock exists, consisting of  $P_1, P_2, P_3$  and  $P_4$ .

	Request		
	A	B	C
P0	0	0	0
P1	2	0	2
P2	0	0	1
P3	1	0	0
P4	0	0	2

# Detection-Algorithm Use

- ❑ When, and how often to invoke depends on:
  - How often a deadlock is likely to occur?
  - How many processes will need to be rolled back?
    - ❑ One for each disjoint cycle
- ❑ How often --
  - Every time a request for allocation cannot be granted immediately
    - ❑ Allows us to detect set of deadlocked processes and process that “caused” deadlock. Extra overhead.
    - ❑ Every hour or whenever CPU utilization drops.
  - With arbitrary invocation there may be many cycles in the resource graph and we would not be able to tell which of the many deadlocked processes “caused” the deadlock.

# Recovery from Deadlock: Process Termination

- ❑ Abort all deadlocked processes.
  - ❑ Abort one process at a time until the deadlock cycle is eliminated.
  - ❑ In which order should we choose to abort?
    - ❑ Priority of the process
    - ❑ How long the process has computed, and how much longer to completion.
    - ❑ Resources the process has used.
    - ❑ Resources process needs to complete.
    - ❑ How many processes will need to be terminated.
    - ❑ Is process interactive or batch?
-

---

# Recovery from Deadlock: Resource Preemption

- Selecting a victim - minimize cost.
  - Rollback
    - return to some safe state, restart process from that state.
  - Starvation
    - same process may always be picked as victim; include number of rollback in cost factor.
-

# Combined approach to deadlock handling

- Combine the three basic approaches

- Prevention
- Avoidance
- Detection

allowing the use of the optimal approach for each class of resources in the system.

- Partition resources into hierarchically ordered classes.

- Use most appropriate technique for handling deadlocks within each class.